

Pranati Modumudi

pranmod01.github.io | (408)-464-8674 | pranati.m@columbia.edu | linkedin.com/in/pranatimodumudi/

EDUCATION

Columbia University

Aug 2025 - May 2027 (Expected)

M.S. in Computer Science, MS Thesis Track (Machine Learning and Computational Neuroscience)

Relevant Coursework: Artificial Intelligence, LLM-Based Generative AI, Computation and the Brain, Theoretical Neuroscience

University of California, Berkeley

Aug 2019 - May 2023

B.A. Data Science Honors and B.A. Economics | Regents and Chancellor's Scholar | Berkeley SCET

Relevant Coursework: Linear Algebra and Differential Equations; Probability & Probabilistic Modeling; Data Engineering; NLP; Game Theory

SKILLS

- ❖ **ML/DL & RL:** PyTorch, JAX, Scikit-learn, Neural Networks, Reinforcement Learning, Model Training & Evaluation, Distributed Training, MLflow, Weights & Biases
 - ❖ **LLMs & Generative AI:** Large Language Models, Prompt Engineering, RAG, RLHF, Fine-Tuning, MCP, Langchain, HuggingFace, AllenNLP, SpaCy
 - ❖ **Programming & Systems:** Python, R, SQL, Bash, Java, Git, Docker, Databricks, Snowflake, Big Data Systems, ETL, APIs
 - ❖ **Neuroscience & Signal Processing:** EEG, Eye-Tracking, Neural Data Analysis, Real-Time Signal Processing, BCIs
-

RESEARCH

Human Decision Modeling | Laboratory for Intelligent Imaging and Neural Computing, Columbia University

- ❖ Developing machine learning pipelines to decode human cognitive states from multi-modal physiological data (EEG, eye-tracking) during decision-making under risk and ambiguity.
- ❖ Built preprocessing and feature extraction pipelines for high-dimensional neural time-series data, implementing classification models for physiological markers of risk aversion and uncertainty.

Undergraduate Senior Honors Thesis | Redwood Center for Theoretical Neuroscience UC Berkeley

- ❖ Modeled hippocampal place cell representations across species (rats → bats), investigating cross-species spatial coding differences advised by Prof. Friedrich Sommer & Dr. Gautam Agarwal.
-

INDUSTRY

Data Scientist I, Disney Streaming, New York

Aug 2023 - May 2025

- ❖ Built production ML pipelines deploying Bayesian structural time-series models and gradient boosted decision trees for causal inference on user behavior, serving 150M+ subscribers with real-time prediction latency <100ms.
- ❖ Developed a distributed training framework for time-series forecasting models (LSTM, Prophet) processing 10M+ daily events, reducing model training time from 6 hours to 45 minutes via parallelized hyperparameter tuning.
- ❖ Engineered reusable Python package for synthetic data generation and continuous treatment effect estimation; deployed via PyPI with automated testing in Jenkins, adopted by five teams for experimental design and model validation.

Machine Learning Systems Intern, OctoML (acquired by Nvidia), Remote

Oct 2022 - Dec 2022

- ❖ Optimized TVM machine learning compiler by running 100+ experiments on CNN models across varying architectures, tuning autoscheduling hyperparameters to reduce convergence time by 10% and improve GPU utilization.
 - ❖ Diagnosed root causes of performance variance in the TVM auto-tuning pipeline through statistical analysis of compilation traces; findings [published](#) in the MLArchSys International Symposium on Computer Architecture 2022.
-

SELECTED PROJECTS

Motor Intent Detection in Hybrid BCIs | Personal Project

- ❖ Built an EEG classification pipeline using CSP+LDA, EEGNet, and LSTM to distinguish volitional motor intent from externally-evoked responses, improving specificity from 51% to 83% for rehabilitation brain-computer interfaces.

Breaking Anti-Sycophancy: Multi-Agent AI Manipulation | Apart Research AI Manipulation Hackathon

- ❖ Designed adversarial framework testing anti-sycophancy defenses across 120 multi-agent deliberations, revealing adversarial pressure achieves 70% manipulation success vs. 15% for sycophantic approaches.

Multi-Agent SLM Evaluators for AI Slop Detection | AWS Small Language Build Day Hackathon

- ❖ Built an ensemble of four specialized small language models (1-3B parameters) fine-tuned with LoRA on AWS Trainium to detect low-fidelity model outputs, enabling order-of-magnitude cost reduction vs. large-model inference for content quality monitoring.